

2/27


- Tested different parameters for the model (convoluted layers, nodes, and dense layers)
- Found the following accuracy for tests:
  - 73% accuracy and 71% validation accuracy for 1 conv-64 nodes-0 dense
  - 0.7220 accuracy and 0.7035 validation accuracy for 1 conv-64 nodes-1 dense
  - 0.7241 accuracy and 0.6935 validation accuracy for 1-conv-128-nodes-0-dense

2/28

- Finished running VGG16 Model
- Accuracy: 0.6703, Loss: 0.6216, Validation Accuracy: 0.6784, Validation Loss: 0.6110


epoch\_acc

epoch\_acc



epoch\_loss


epoch\_loss



.

epoch\_val\_acc

epoch\_val\_acc



Full output of VGG16:

VGG16

WARNING:tensorflow:From

/usr/local/lib/python3.6/dist-packages/tensorflow\_core/python/ops/resource\_variable\_ops.py:1630: calling BaseResourceVariable.\_\_init\_\_ (from tensorflow.python.ops.resource\_variable\_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass \*\_constraint arguments to layers.

Downloading data from

[https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels.h5](https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_kernels.h5)

553467904/553467096 [=====] - 34s 0us/step

Train on 464 samples, validate on 199 samples

Epoch 1/50

464/464 [=====] - 342s 736ms/sample - loss: 0.7849 - acc: 0.5625 - val\_loss: 0.6902 - val\_acc: 0.5829

Epoch 2/50

464/464 [=====] - 341s 736ms/sample - loss: 0.7076 - acc:  
0.5259 - val\_loss: 0.6673 - val\_acc: 0.5930  
Epoch 3/50  
464/464 [=====] - 340s 733ms/sample - loss: 0.6967 - acc:  
0.5453 - val\_loss: 0.6624 - val\_acc: 0.5930  
Epoch 4/50  
464/464 [=====] - 339s 730ms/sample - loss: 0.6875 - acc:  
0.5647 - val\_loss: 0.6586 - val\_acc: 0.6382  
Epoch 5/50  
464/464 [=====] - 340s 732ms/sample - loss: 0.6800 - acc:  
0.5862 - val\_loss: 0.6565 - val\_acc: 0.6533  
Epoch 6/50  
464/464 [=====] - 338s 729ms/sample - loss: 0.6759 - acc:  
0.6121 - val\_loss: 0.6542 - val\_acc: 0.6482  
Epoch 7/50  
464/464 [=====] - 338s 728ms/sample - loss: 0.6732 - acc:  
0.6142 - val\_loss: 0.6519 - val\_acc: 0.6382  
Epoch 8/50  
464/464 [=====] - 338s 729ms/sample - loss: 0.6706 - acc:  
0.6315 - val\_loss: 0.6496 - val\_acc: 0.6382  
Epoch 9/50  
464/464 [=====] - 338s 729ms/sample - loss: 0.6679 - acc:  
0.6336 - val\_loss: 0.6474 - val\_acc: 0.6533  
Epoch 10/50  
464/464 [=====] - 338s 729ms/sample - loss: 0.6654 - acc:  
0.6293 - val\_loss: 0.6453 - val\_acc: 0.6583  
Epoch 11/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6631 - acc:  
0.6422 - val\_loss: 0.6434 - val\_acc: 0.6683  
Epoch 12/50  
464/464 [=====] - 336s 725ms/sample - loss: 0.6609 - acc:  
0.6422 - val\_loss: 0.6415 - val\_acc: 0.6683  
Epoch 13/50  
464/464 [=====] - 338s 728ms/sample - loss: 0.6588 - acc:  
0.6466 - val\_loss: 0.6398 - val\_acc: 0.6633  
Epoch 14/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6568 - acc:  
0.6487 - val\_loss: 0.6382 - val\_acc: 0.6533  
Epoch 15/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6550 - acc:  
0.6444 - val\_loss: 0.6366 - val\_acc: 0.6583  
Epoch 16/50

464/464 [=====] - 337s 726ms/sample - loss: 0.6532 - acc:  
0.6509 - val\_loss: 0.6352 - val\_acc: 0.6533  
Epoch 17/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6516 - acc:  
0.6552 - val\_loss: 0.6338 - val\_acc: 0.6533  
Epoch 18/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6500 - acc:  
0.6552 - val\_loss: 0.6326 - val\_acc: 0.6432  
Epoch 19/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6485 - acc:  
0.6573 - val\_loss: 0.6313 - val\_acc: 0.6533  
Epoch 20/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6470 - acc:  
0.6616 - val\_loss: 0.6302 - val\_acc: 0.6533  
Epoch 21/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6457 - acc:  
0.6638 - val\_loss: 0.6291 - val\_acc: 0.6533  
Epoch 22/50  
464/464 [=====] - 337s 725ms/sample - loss: 0.6443 - acc:  
0.6659 - val\_loss: 0.6281 - val\_acc: 0.6633  
Epoch 23/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6431 - acc:  
0.6659 - val\_loss: 0.6271 - val\_acc: 0.6633  
Epoch 24/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6419 - acc:  
0.6659 - val\_loss: 0.6262 - val\_acc: 0.6683  
Epoch 25/50  
464/464 [=====] - 338s 729ms/sample - loss: 0.6408 - acc:  
0.6659 - val\_loss: 0.6253 - val\_acc: 0.6683  
Epoch 26/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6397 - acc:  
0.6659 - val\_loss: 0.6244 - val\_acc: 0.6683  
Epoch 27/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6386 - acc:  
0.6703 - val\_loss: 0.6236 - val\_acc: 0.6683  
Epoch 28/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6376 - acc:  
0.6681 - val\_loss: 0.6229 - val\_acc: 0.6683  
Epoch 29/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6366 - acc:  
0.6681 - val\_loss: 0.6221 - val\_acc: 0.6683  
Epoch 30/50


464/464 [=====] - 337s 727ms/sample - loss: 0.6357 - acc:  
0.6659 - val\_loss: 0.6214 - val\_acc: 0.6784  
Epoch 31/50  
464/464 [=====] - 338s 728ms/sample - loss: 0.6348 - acc:  
0.6681 - val\_loss: 0.6207 - val\_acc: 0.6734  
Epoch 32/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6339 - acc:  
0.6681 - val\_loss: 0.6201 - val\_acc: 0.6734  
Epoch 33/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6330 - acc:  
0.6703 - val\_loss: 0.6194 - val\_acc: 0.6784  
Epoch 34/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6322 - acc:  
0.6681 - val\_loss: 0.6188 - val\_acc: 0.6784  
Epoch 35/50  
464/464 [=====] - 336s 725ms/sample - loss: 0.6314 - acc:  
0.6659 - val\_loss: 0.6182 - val\_acc: 0.6784  
Epoch 36/50  
464/464 [=====] - 336s 725ms/sample - loss: 0.6306 - acc:  
0.6681 - val\_loss: 0.6176 - val\_acc: 0.6734  
Epoch 37/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6299 - acc:  
0.6703 - val\_loss: 0.6171 - val\_acc: 0.6734  
Epoch 38/50  
464/464 [=====] - 336s 725ms/sample - loss: 0.6291 - acc:  
0.6724 - val\_loss: 0.6165 - val\_acc: 0.6734  
Epoch 39/50  
464/464 [=====] - 337s 725ms/sample - loss: 0.6284 - acc:  
0.6746 - val\_loss: 0.6160 - val\_acc: 0.6734  
Epoch 40/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6277 - acc:  
0.6746 - val\_loss: 0.6155 - val\_acc: 0.6784  
Epoch 41/50  
464/464 [=====] - 338s 728ms/sample - loss: 0.6271 - acc:  
0.6724 - val\_loss: 0.6150 - val\_acc: 0.6784  
Epoch 42/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6264 - acc:  
0.6746 - val\_loss: 0.6145 - val\_acc: 0.6734  
Epoch 43/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6257 - acc:  
0.6724 - val\_loss: 0.6140 - val\_acc: 0.6734  
Epoch 44/50

464/464 [=====] - 337s 726ms/sample - loss: 0.6251 - acc:  
0.6724 - val\_loss: 0.6136 - val\_acc: 0.6734  
Epoch 45/50  
464/464 [=====] - 337s 726ms/sample - loss: 0.6245 - acc:  
0.6724 - val\_loss: 0.6131 - val\_acc: 0.6734  
Epoch 46/50  
464/464 [=====] - 338s 728ms/sample - loss: 0.6239 - acc:  
0.6724 - val\_loss: 0.6127 - val\_acc: 0.6784  
Epoch 47/50  
464/464 [=====] - 337s 727ms/sample - loss: 0.6233 - acc:  
0.6703 - val\_loss: 0.6122 - val\_acc: 0.6784  
Epoch 48/50  
464/464 [=====] - 336s 725ms/sample - loss: 0.6227 - acc:  
0.6703 - val\_loss: 0.6118 - val\_acc: 0.6784  
Epoch 49/50  
464/464 [=====] - 336s 724ms/sample - loss: 0.6221 - acc:  
0.6703 - val\_loss: 0.6114 - val\_acc: 0.6784  
Epoch 50/50  
464/464 [=====] - 336s 723ms/sample - loss: 0.6216 - acc:  
0.6703 - val\_loss: 0.6110 - val\_acc: 0.6784  
<tensorflow.python.keras.callbacks.History at 0x7f4fe6a7f860>

epoch\_val\_loss

---

epoch\_val\_loss




- 
- Retrained CNN model
  - Had some issues with Tensorboard loading data (scalars) as graphs, so copied code into a new file and ran again, which worked



epoch\_acc

epoch\_acc




•

epoch\_loss

---


epoch\_loss



.

epoch\_val\_acc


epoch\_val\_acc



•

## epoch\_val\_loss

epoch\_val\_loss



- 

- Full output:

1-conv-64-nodes-0-dense-1582954102

WARNING:tensorflow:From

/usr/local/lib/python3.6/dist-packages/tensorflow\_core/python/ops/resource\_variable\_ops.py:16

30: calling BaseResourceVariable.\_\_init\_\_ (from tensorflow.python.ops.resource\_variable\_ops)

with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass \*\_constraint arguments to layers.

WARNING:tensorflow:From

/usr/local/lib/python3.6/dist-packages/tensorflow\_core/python/ops/nn\_impl.py:183: where (from

tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 464 samples, validate on 199 samples

Epoch 1/10

464/464 [=====] - 37s 79ms/sample - loss: 3.2037 - acc:

0.5323 - val\_loss: 2.0608 - val\_acc: 0.4774

Epoch 2/10

464/464 [=====] - 33s 71ms/sample - loss: 0.9785 - acc:  
0.4892 - val\_loss: 0.6749 - val\_acc: 0.5829  
Epoch 3/10  
464/464 [=====] - 32s 70ms/sample - loss: 0.6615 - acc:  
0.6099 - val\_loss: 0.6781 - val\_acc: 0.5930  
Epoch 4/10  
464/464 [=====] - 33s 70ms/sample - loss: 0.6368 - acc:  
0.6293 - val\_loss: 0.6660 - val\_acc: 0.5578  
Epoch 5/10  
464/464 [=====] - 33s 71ms/sample - loss: 0.6133 - acc:  
0.6595 - val\_loss: 0.6628 - val\_acc: 0.5628  
Epoch 6/10  
464/464 [=====] - 33s 71ms/sample - loss: 0.5794 - acc:  
0.7026 - val\_loss: 0.6707 - val\_acc: 0.5678  
Epoch 7/10  
464/464 [=====] - 33s 70ms/sample - loss: 0.5807 - acc:  
0.6940 - val\_loss: 0.6696 - val\_acc: 0.5779  
Epoch 8/10  
464/464 [=====] - 33s 71ms/sample - loss: 0.5542 - acc:  
0.7134 - val\_loss: 0.6544 - val\_acc: 0.6231  
Epoch 9/10  
464/464 [=====] - 33s 71ms/sample - loss: 0.5415 - acc:  
0.7435 - val\_loss: 0.6739 - val\_acc: 0.5879  
Epoch 10/10  
464/464 [=====] - 33s 72ms/sample - loss: 0.5115 - acc:  
0.7478 - val\_loss: 0.6519 - val\_acc: 0.5930